

— COGNITIONX EGYPT · 3 AUGUST 2026

AI-Native Development

From Writing Code to Owning the Full SDLC

A role-evolution talk, not a tools talk.

Direction · Constraints · Evidence · Outcomes

Hany Saad Senior Engineering Manager, ITWorx



Connect on LinkedIn

You didn't write this code.

But your name is on the release.

RELEASE CANDIDATE RC-2026.08		PREPARED FOR PRODUCTION	
3,247 AI-generated lines	<i>Would you ship it?</i>		PASSED
			PASSED
41 files changed	⚠ Architecture boundary changed		REVIEW
	⚠ New dependency introduced		REVIEW
	❗ Security review incomplete		INCOMPLETE

AI can generate the implementation.

The engineer still owns the outcome.

The Bottleneck Moved

Generating code

is becoming easier.



**Knowing whether it deserves to
ship**

is becoming harder.

The new bottleneck is **verification**.

Do not surrender control to AI.
Move your control to a higher level.



You are the orchestrator—not the typist.

AI in Development Is No Longer Optional

Use the right level of AI for the right task.



01

Vibe Coding

BEST FOR

Rapid prototypes, MVP experiments, demos, hackathons, idea validation, and disposable spikes.

Optimize for learning speed—not production confidence.

EXAMPLES

Lovable

bolt.new

Google AI Studio



02

AI-Assisted Coding

BEST FOR

Daily coding, debugging, refactoring, tests, documentation, explanation, and unfamiliar code.

AI proposes; the engineer reads, integrates, and remains responsible.

EXAMPLES

GitHub Copilot

Cursor

Equivalent IDE assistants



03

AI Coding Agents

BEST FOR

Well-specified issues, repository analysis, test repair, controlled modernization, and PR preparation.

Delegate the task—not the accountability.

EXAMPLES

Claude Code

OpenAI Codex

Gemini CLI

Copilot coding agent

The question is no longer whether to use AI—but how to use it effectively, safely, and responsibly.

Tool examples illustrate categories and are not endorsements. Choose tools according to task, security, privacy, cost, and organizational policy.

— SAME PROMISE. DIFFERENT CONDITIONS.

+55.8%

FASTER

Controlled, self-contained task

Peng et al. • Microsoft Research / arXiv • 2023

-19%


SLOWER

Experienced developers, mature
repositories

METR • Mature repository study • 2025

Both results are real. What engineering conditions explain the difference?

Prompt → Context → Harness → Loop



- | | | |
|---|----------------|--|
| 1 | Prompt | What are we asking for? |
| 2 | Context | What does the agent need to know? |
| 3 | Harness | What can it access, execute, and change? |
| 4 | Loop | How does it verify, retry, stop, and escalate? |

Each layer contains—and depends on—the one before it.

— RESPONSIBILITY EXPANDS WITH LEVERAGE


Three Loops. Three Clocks.



The inner loop got fast. Your job moved outward.

From Sequential Delivery to a Continuous, Human-Guided Loop

The lifecycle remains. Execution accelerates. Human judgment moves to the control points.



TRADITIONAL DELIVERY

Sequential handoffs

- Humans create most artifacts manually
- Coding consumes much of the delivery cycle
- Testing and review follow implementation
- Documentation and operational readiness often lag
- Feedback arrives after release



AI-NATIVE DELIVERY

Continuous, human-guided loops

- Humans frame, specify, constrain, and approve
- AI accelerates bounded execution
- Tests and evidence are defined earlier
- Documentation and delivery artifacts evolve with the change
- Production feedback continuously updates the next cycle



The agent accelerates the work inside the lifecycle. The engineer still owns the lifecycle.

What Owning the Full SDLC Actually Means

- | | | |
|----|----------------------------|---|
| 01 | Intent | Problem, outcome, and the decision not to build |
| 02 | Specification | Acceptance criteria and edge cases |
| 03 | Architecture | Boundaries, contracts, and data flows |
| 04 | Delegation boundary | Scope, access, tools, and stopping conditions |
| 05 | Evidence pack | Tests, security, performance, and observability |
| 06 | Release case | Rollout, provenance, and rollback |
| 07 | Production learning | Telemetry, incidents, feedback, and improvement |

Every artifact has an owner. AI is not that owner.

Start With Intent, Not Code

“Prevent session overbooking. If capacity is reached, offer a waitlist and notify organizers.”

Raw business request

What counts as capacity?

Who may override it?

What personal data is exposed?

What is the rollback behavior?

Is check-in concurrent?

How is the waitlist ordered?

What happens when notification fails?

ACCEPTANCE CRITERIA

Atomic capacity check

Authorized override

Stable waitlist order

Failure-tolerant notification

Auditable rollback

The engineer creates value before the first line is generated.

— CHALLENGE THE PLAN BEFORE ACCEPTING THE WORK

Plan. Constrain. Delegate.

AGENT'S PROPOSED PLAN

- 01 Add waitlist storage
- 02 Update check-in endpoint
- 03 Add organizer notification
- 04 Update attendee UI
- 05 Add tests


HUMAN REVIEW

- Concurrency is missing
- Authorization boundary is unclear
- Notification failure must not reverse check-in
- No new dependency without approval
- Audit event required

Delegation contract bounded tasks · affected interfaces · tests required · permissions allowed · stop / escalation conditions

Verify the System, Not the Diff

Verify the behavior, boundaries, operations, and critical decisions—not only the generated code.



Green tests show execution. Change assurance shows the requirement survived.

Code Review Is Necessary. It Is No Longer Sufficient.

01 Business intent

02 Acceptance behavior

03 Interfaces and data contracts

04 Architecture and dependencies

05 Security and supply chain

06 Operability and rollback

07 Critical code

Do not stop reviewing code.

Stop pretending code review alone proves the system is safe.

The Engineer Moves Up the Stack

AI fluency builds on engineering maturity—not instead of it.

01	 Software Fundamentals	Programming · OOP · data structures · APIs · databases · debugging · Git
02	 Engineering Practice	Testing · code review · CI/CD · secure coding · observability
03	 System Design	Architecture · boundaries · data contracts · reliability · trade-offs
04	 Product & Domain Judgment	User needs · business rules · acceptance criteria · domain context
05	 AI Collaboration	Task framing · context · constraints · tool use · critique
06	 End-to-End Ownership	Verification · risk · orchestration · release decisions · outcomes

HOW CAPABILITY COMPOUNDS

- 01-02
Foundation
Build and ship sound software.
- 03-04
Maturity
Shape systems and problems.
- 05-06
Leverage
Direct, verify, and own.

AI amplifies the stack.
It does not replace its foundation.

Your AI-Native Starting Playbook

1

Choose one bounded, real task
Useful enough to matter. Small enough to verify.

2

Define acceptance criteria first
Expected behavior, edge cases, constraints, and evidence.

3

Make repository context explicit
Architecture, contracts, conventions, security, and Definition of Done.

4

Put automated evidence before human attention
Run tests, scans, policy checks, and contract validation first.

5

Measure delivery—not generated lines
Track outcomes, review effort, rework, defects, and release stability.

CHOOSE THE RIGHT OPERATING MODE



Increase autonomy only when specification, context, and verification are ready.

MAKE THE REPOSITORY AGENT-READY

DIRECTION	AGENTS.md	CLAUDE.md	.github/copilot-instructions.md
	Architecture decisions	Definition of Done	
BOUNDARIES	Security policies	API contracts	
EVIDENCE	Test strategy	Acceptance criteria	

MEASURE OUTCOMES — NOT GENERATED LINES				
Cycle time	Review effort	Rework	Escaped defects	Change-failure rate

A Project Structure for Coding Agents

Turn repeated prompts into versioned, reviewable operating assets.

product-service/ SAMPLE – ADAPT PATHS TO YOUR RUNTIME

INSTRUCTIONS	AGENTS.md .github/copilot-instructions.md .github/instructions/*.instructions.md	Always-on and path-specific rules
AGENTS	.github/agents/ planner.md implementer.md reviewer.md	Roles, tools, and delegated scope
SKILLS	.github/skills/ testing/SKILL.md security/SKILL.md release/SKILL.md	Reusable workflows and resources
HOOKS	.github/hooks/ policy.json validation.json	Deterministic lifecycle checks
TOOLS + GATES	.github/mcp.json scripts/agent/verify.ps1	Controlled tools and evidence entry points

BEFORE WORK BEGINS

The repository should answer:

01 Which instructions apply?

02 Which agent owns the task?

03 Which skills should load?

04 Which hooks enforce the gate?

Prompts initiate work. Repository configuration makes it repeatable.

— OWN THE OUTER LOOPS

You are the orchestrator—not the typist.

The inner loop got fast. Own the outer loops.


AI-native development is not about producing more code.

It is about owning the system that turns intent into reliable outcomes.



Hany Saad

Senior Engineering Manager, ITWorx



LinkedIn



Presentation repo